

BLADE: Bilevel Low-rank Augmented-Lagrangian Driven Erasure for LLM Unlearning

Anonymous ACL submission

Abstract

Existing LLM unlearning methods struggle with robustness: unbounded forget losses degrade model coherence, fixed-weight balancing cannot adapt as retain difficulty shifts mid-training, and methods that work on one benchmark falter under scaling or repeated application. We propose BLADE, a constrained bilevel framework that gives smooth, predictable control over the optimization landscape. The forget objective is *clamped entropy*—per-token entropy maximized only to a vocabulary-invariant threshold $\tau \cdot \log V$, after which the gradient is exactly zero—yielding a bounded, self-decelerating loss. Retain quality is enforced as a hard inequality constraint via an augmented Lagrangian with asymmetric dual updates, nested in a bilevel structure where the inner loop repairs retain damage before each outer forgetting step; all updates are confined to LoRA adapters. BLADE dominates across three benchmark families—improving composite scores over the strongest baselines by 8% on TOFU, 9% on MUSE Books, and 7% on KnowUndo—and remains stable under $4\times$ scaling and 4 sequential unlearning steps on MUSE News where the best competing method collapses entirely.

1 Introduction

Deployed language models require continuous post-hoc correction: regulations such as the GDPR (Voigt and von dem Bussche, 2017) mandate deletion of personal data on request, copyright holders demand removal of memorized content (Carlini et al., 2021), and models that absorb misinformation—recently demonstrated when LLMs confidently repeated a fabricated disease (Stokel-Walker, 2026)—must be patched without full retraining. *Machine unlearning* addresses these needs, yet the pro-

cess is inherently *recurring*: new removal requests, policy changes, and safety incidents arrive throughout a model’s lifetime (Liu et al., 2025).

This makes the central challenge not just *what* to forget but *how to forget without breaking what remains*. Forget and retain knowledge share entangled representations (Cheng et al., 2026), so each unlearning step risks collateral damage that compounds over repeated applications. Existing methods rely on fixed-weight loss balancing (Liu et al., 2025; Zhang et al., 2024) that cannot adapt as retain difficulty shifts mid-training, producing three recurring failure modes: (i) *unbounded forgetting*—gradient ascent (Jang et al., 2023) produces arbitrarily large gradients, destroying utility; (ii) *static tradeoff*—fixed-weight methods (GradDiff, NPO) oscillate without self-correction; (iii) *over-forgetting*—unclamped entropy (Yuan et al., 2025) pushes already-forgotten tokens, wasting gradient budget on tokens that no longer need perturbation.

BLADE addresses these through a formulation designed for *robust, smooth, and manageable* unlearning. A *clamped entropy* forget loss (Section 3.4) maximizes per-token entropy only up to a vocabulary-invariant threshold $\tau \cdot \log V$, after which the gradient is exactly zero—bounded, self-decelerating, and reference-free. Retain quality is enforced as an inequality constraint via an augmented Lagrangian with asymmetric dual updates (Section 3.5) that permanently ratchet protection after any violation, nested in a bilevel structure (Section 3.3) whose inner loop repairs retain damage before each outer forgetting step; all updates are confined to LoRA adapters (Section 3.2), bounding representational drift.

We evaluate on TOFU (Maini et al., 2024) (Llama-3.2-1B/3B, three splits, 5 seeds),

MUSE (Shi et al., 2025) (Books/News, Llama-2-7B), and KnowUndo (Tian et al., 2024) (copyright/privacy, Llama-2-7B-chat). Key findings:

- **Dominant across benchmarks.** BLADE improves composite scores over the strongest baselines by 8% on TOFU, 9% on MUSE Books, and 7% on KnowUndo, with near-zero semantic leakage confirmed by an LLM judge.
- **Robust under stress.** Under $4\times$ scaling and 4 sequential unlearning steps, BLADE maintains stable performance (HM 0.533–0.552) while the strongest baseline collapses catastrophically (HM 0.58→0.005).
- **Smooth, controllable optimization.** Training consistently follows a three-phase pattern—warm-up, spike-and-ratchet, smooth convergence—giving interpretable control over the forget–retain tradeoff absent from all baselines, whose dynamics oscillate without self-correction.

2 Related Work

LLM unlearning methods. Machine unlearning originated in the classification setting (Cao and Yang, 2015; Bourtole et al., 2019); see Liu et al. (2024); Blanco-Justicia et al. (2024) for surveys. Early approaches used Fisher information to scrub data influence (Golatkhar et al., 2020a,b) and influence functions (Koh and Liang, 2017) to identify responsible training points—later extended to iterative LLM unlearning by SOUL (Jia et al., 2024b)—but these do not transfer directly to autoregressive LLMs at scale. For LLMs, gradient ascent (GA) (Jang et al., 2023) maximizes the forget loss but diverges and destroys coherence; GradDiff (Yao et al., 2024) adds retain cross-entropy but still degrades; and Eldan and Russinovich (2023) demonstrated approximate unlearning via label replacement. Preference-based methods—NPO (Zhang et al., 2024), SimNPO (Fan et al., 2025), NGDiff (Jin et al., 2025)—collapse more gracefully but saturate without converging to a well-defined target. RMU (Li et al., 2024) steers activations toward random targets for forget data. What these approaches share is the absence of a *bounded* forget loss with a per-token stopping criterion: GA diverges, NPO damps but never

reaches zero, and unclamped entropy (Yuan et al., 2025) pushes indefinitely.

Parameter-efficient unlearning. LoRA (Hu et al., 2022) has been applied to unlearning via Fisher-informed initialization (LoKU; Cha et al. 2025) and importance-weighted masking (Jia et al., 2024a). In BLADE, LoRA serves as a *structural constraint*: the low-rank bottleneck itself prevents catastrophic drift, analogous to how EWC (Kirkpatrick et al., 2016) protects important weights in continual learning but applied structurally rather than via per-weight penalties.

Benchmarks. TOFU (Maini et al., 2024) tests factual associations with synthetic biographies and provides a gold retrained model for reference. MUSE (Shi et al., 2025) tests naturalistic memorization at scale with explicit scalability and sustainability stress tests. KnowUndo (Tian et al., 2024) targets domain-specific unlearning (copyright, privacy) from LoRA-fine-tuned models. We evaluate on all three for complementary coverage.

Bilevel and constrained optimization. BLUR (Reisizadeh et al., 2026) places forget in the lower level and retain in the upper, resolving conflicts via gradient projection. PDU (Entesari et al., 2026) uses a linear Lagrangian with logit-margin loss and symmetric dual updates. Cheng et al. (2026) apply ALM with an equality constraint on vision models only. BLADE differs: (a) an *inequality* constraint that never penalizes retain below ε ; (b) *asymmetric* dual updates producing ratchet behavior; (c) retain in the *inner loop* (repair-first); and (d) clamped entropy as the forget loss—bounded with a hard per-token deadzone.

3 Method

Unlearning often degrades general capability because forget and retain knowledge occupy shared representations (Cheng et al., 2026; Liu et al., 2025). BLADE addresses this with three mechanisms: a bilevel structure that repairs retain *before* each forget step (§3.3), a clamped entropy loss with a hard per-token deadzone (§3.4), and an augmented Lagrangian whose dual variable adapts the forget–retain tradeoff during training (§3.5).

3.1 Problem Setup

Let M_θ denote an LLM parameterized by θ (initialized from pretrained weights θ_0), $\mathcal{D}_{\text{forget}}$ the data to unlearn, and $\mathcal{D}_{\text{retain}}$ data the model should preserve. We seek θ such that (i) the output distribution of M_θ on $\mathcal{D}_{\text{forget}}$ is high-entropy (memorized content is forgotten) and (ii) the cross-entropy loss on $\mathcal{D}_{\text{retain}}$ remains below a threshold ε .

3.2 LoRA Parameterization

Rather than modifying full weight matrices, we attach LoRA adapters (Hu et al., 2022) to all projection matrices ($q/k/v/o_proj$, $gate/up/down_proj$):

$$W' = W_0 + \frac{\alpha}{r} BA, \quad B \in \mathbb{R}^{d \times r}, \quad A \in \mathbb{R}^{r \times d} \quad (1)$$

where $r \ll d$ is the rank and α/r is a scaling factor. With $r=16$ on a 7B model, only $\sim 0.2\%$ of parameters are trainable. The low-rank bottleneck limits what the update can express, preventing catastrophic weight drift even under aggressive forgetting gradients. Another practical benefit is that disabling the adapters keeps the original model intact which simplifies rollback, debugging and subsequent forget/retain requests. With the parameter space constrained, the remaining question is how to organize the optimization itself.

3.3 Bilevel Formulation

We decompose the problem into two nested optimization levels: *retain in the inner loop, forget in the outer loop*.

Inner loop (retain preservation). For each outer step, we run K SGD steps on the retain set:

$$\theta \leftarrow \theta - \eta_{\text{in}} \nabla_{\theta} \mathcal{L}_{\text{CE}}(M_{\theta}; \mathcal{B}_{\text{ret}}) \quad (2)$$

where $\mathcal{B}_{\text{ret}} \sim \mathcal{D}_{\text{retain}}$ is a mini-batch and $\mathcal{L}_{\text{CE}}$ is cross-entropy. The inner loop uses a reliably faster learning rate (η_{in}) to restore any retain damage from the previous outer step.

Outer loop (constrained forgetting).

The outer objective combines the forget loss with an inequality constraint on retain performance, enforced via an augmented Lagrangian:

$$\mathcal{L}_{\text{outer}} = \mathcal{L}_{\text{fgt}} + \lambda r + \frac{\rho}{2} [\max(0, r)]^2, \quad r = L_{\text{ret}} - \varepsilon \quad (3)$$

where $L_{\text{ret}} = \mathcal{L}_{\text{CE}}(M_{\theta}; \mathcal{B}'_{\text{ret}})$ is evaluated on a *fresh* retain batch (separate from the inner loop), λ is the Lagrange multiplier, ρ is the quadratic penalty coefficient, and ε is the retain budget. The outer loop uses a slower learning rate (η_{out}) compared to the inner loop, keeping each forgetting step small enough for the inner loop to repair.

Repair-first principle. Prior two-level unlearning methods like SCRUB (Kurmanji et al., 2023) (min-max) and BLUR (Reisizadeh et al., 2026) (bilevel with gradient projection) prioritize forgetting at the lower level and treat retain as the upper-level concern. We invert this assignment.

Forgetting is inherently destructive: gradient ascent, entropy maximization, and preference-based losses all push the model *away* from learned representations, causing catastrophic collapse when unconstrained (Zhang et al., 2024). To make matters worse, forget and retain knowledge share entangled representations (Cheng et al., 2026), so perturbing forget tokens inevitably disturbs retain performance.

Retain repair therefore belongs in the inner loop as a stabilizing anchor. It runs first and comparably faster ($\eta_{\text{in}} > \eta_{\text{out}}$), so the model enters each outer step from a state with good retain performance. The outer loop then applies one small forgetting step, sized so the next inner loop can fully repair any collateral damage. This prevents the accumulation of representational drift that makes recovery progressively harder when the loops are reversed.

For this structure to work, the forget loss in the outer loop must be well-behaved: bounded, self-saturating, and compatible with small step sizes. We introduce *clamped entropy* to address this.

3.4 Clamped Entropy Loss

Entropy maximization as a forget objective is well-known (Yuan et al., 2025), but unclamped formulations apply nonzero gradients to every token at every step, even those already near-uniform—needlessly perturbing shared representations. We introduce *clamped entropy*, which defines when a token is “sufficiently for-

gotten” and stops optimizing it:

$$\mathcal{L}_{\text{forget}} = \frac{1}{N} \sum_{t=1}^N \max(0, \tau \cdot H_{\text{max}} - H(p_t)) \quad (4)$$

where $p_t = \text{softmax}(z_t)$ is the output distribution at position t , $H(p_t) = -\sum_v p_t(v) \log p_t(v)$ is its Shannon entropy, $H_{\text{max}} = \log V$ is the maximum entropy over vocabulary V , and $\tau \in (0, 1)$ is the target fraction. Normalizing by H_{max} makes τ vocabulary-invariant: the same τ applies regardless of whether the vocabulary has 32K (Touvron et al., 2023) or 128K tokens (Grattafiori et al., 2024).

The $\max(0, \cdot)$ clamp produces a hard deadzone: once $H(p_t) \geq \tau \cdot H_{\text{max}}$, token t ’s gradient is *exactly zero*. This makes the loss bounded ($\mathcal{L}_{\text{forget}} \in [0, \tau \cdot H_{\text{max}}]$), self-stabilizing (the effective perturbation shrinks as tokens are forgotten), and reference-free (no frozen model needed, unlike NPO). The bounded outer perturbation (Proposition 1) ensures the inner loop can reliably repair retain damage at any training stage.

Clamped entropy yields well-behaved forget gradients, but the remaining challenge is adapting *how much* retain protection the outer loop enforces as the loss landscape evolves.

3.5 Augmented Lagrangian with Asymmetric Dual Update

The multiplier λ controls how aggressively the outer objective protects retain performance. Rather than treating it as a fixed hyperparameter, we learn it via a dual update rule.

Standard ALM uses symmetric updates designed for equality constraints (Bertsekas, 1982; Nocedal and Wright, 2006). We instead enforce the *inequality* $L_{\text{ret}} \leq \varepsilon$: the quadratic penalty $\frac{\rho}{2}[\max(0, r)]^2$ activates only when violated (Eq. 3), so the model is never penalized for retain loss below ε . On top of this, we apply an *asymmetric* dual update:

$$\lambda \leftarrow \begin{cases} \lambda + \rho |r| & \text{if } r > 0 \\ \lambda - \alpha \rho |r| & \text{if } r \leq 0 \end{cases} \quad (5)$$

where $\alpha \ll 1$. The slower decay means that once a retain violation drives λ up, it stays elevated—preventing the repeated oscillation cycles observed with symmetric updates (PDU (Entesari et al., 2026); see Figure 4).

Algorithm 1 BLADE: Bilevel Augmented Lagrangian Unlearning

Require: Model M with frozen weights θ_0 , $\mathcal{D}_{\text{forget}}$, $\mathcal{D}_{\text{retain}}$

Require: Hyperparams: K , η_{in} , η_{out} , ε_{mul} , ρ , λ_0 , α , τ , T

- 1: Attach LoRA adapters to M ; $\theta \leftarrow$ LoRA params
- 2: $\lambda \leftarrow \lambda_0$
- 3: $\varepsilon \leftarrow \varepsilon_{\text{mul}} \cdot \frac{1}{K} \sum_{k=1}^K \mathcal{L}_{\text{CE}}^{(k)}(\mathcal{D}_{\text{retain}})$
- 4: **for** $t = 1$ **to** T **do**
- 5: **Inner loop** (retain repair):
- 6: **for** $k = 1$ **to** K **do**
- 7: $\theta \leftarrow \theta - \eta_{\text{in}} \nabla_{\theta} \mathcal{L}_{\text{CE}}(M\theta; \mathcal{B}_{\text{ret}})$
- 8: **end for**
- 9: **Forget loss** (clamped entropy):
- 10: $\mathcal{L}_{\text{fgt}} \leftarrow \frac{1}{N} \sum_{i=1}^N \max(0, \tau \log V - H(p_i))$
- 11: **Outer loss** (augmented Lagrangian):
- 12: $r \leftarrow \mathcal{L}_{\text{CE}}(M\theta; \mathcal{B}_{\text{ret}}) - \varepsilon$
- 13: $\mathcal{L}_{\text{outer}} \leftarrow \mathcal{L}_{\text{fgt}} + \lambda r + \frac{\rho}{2}[\max(0, r)]^2$
- 14: $\theta \leftarrow \theta - \eta_{\text{out}} \cdot \text{Adam}(\nabla_{\theta} \mathcal{L}_{\text{outer}})$
- 15: **Dual update** (asymmetric ratchet):
- 16: **if** $r > 0$ **then**
- 17: $\lambda \leftarrow \lambda + \rho |r|$
- 18: **else**
- 19: $\lambda \leftarrow \lambda - \alpha \rho |r|$
- 20: **end if**
- 21: **end for**
- 22: **return** M with trained LoRA adapters

Rather than hand-tuning ε , we set it as a fraction (ε_{mul}) of the model’s average retain loss over one inner loop, so the constraint budget scales with the model’s own retain difficulty and requires no manual calibration across benchmarks (Appendix D).

Convergence criterion. Training terminates when an exponential moving average of the per-step $\mathcal{L}_{\text{fgt}}$ change drops below a fraction of its peak—indicating that remaining tokens have entered the deadzone. A safety cap T prevents runaway training (Appendix B.3).

3.6 Algorithm Summary

Algorithm 1 summarizes the complete procedure: each outer step first runs K inner retain-repair steps, then computes the augmented Lagrangian outer loss, and finally updates the dual variable asymmetrically.

4 Experiments

We evaluate BLADE on three benchmarks that test complementary aspects of unlearning: factual association removal, naturalistic memorization at scale, and domain-specific knowledge extraction from adapted models. All methods are implemented and evaluated within the OpenUnlearning framework (Shi et al., 2025) to ensure fair comparison under

identical data processing, evaluation protocols, and compute budgets. Baseline hyperparameters are listed in Appendix Table 6. All experiments report means and standard deviations over 5 random seeds unless noted.

4.1 Setup

Benchmarks. TOFU (Maini et al., 2024) provides a controlled testbed with 200 synthetic author biographies and QA pairs, evaluated on Llama-3.2-1B/3B-Instruct across three forget splits (1%/5%/10%). Each split removes an increasing fraction of the training authors, testing whether the method scales gracefully. A gold retrained model provides an upper bound.

MUSE (Shi et al., 2025) tests naturalistic memorization on Llama-2-7B across two corpora: *Books* (Harry Potter, a single cohesive document) and *News* (889 BBC articles, a distributed multi-document set). Beyond standard forget/retain performance, MUSE also tests scalability (larger forget sets) and sustainability (sequential unlearning from the same checkpoint), which we evaluate in §4.6.

KnowUndo (Tian et al., 2024) targets copyright and privacy domains on Llama-2-7B-chat that was LoRA-fine-tuned on domain-specific corpora. Unlike MUSE where memorization arises from pretraining, here the target knowledge was explicitly injected via LoRA adaptation—testing whether unlearning can reverse fine-tuning without destroying general capability.

Metrics. Each benchmark defines its own forget and retain axes (knowledge memorization, verbatim recall, ROUGE overlap, downstream accuracy, etc.). To enable cross-benchmark comparison, we report a *harmonic mean* (HM) that aggregates the key forget and retain metrics for each dataset, penalizing methods that sacrifice one axis for another. Detailed per-benchmark HM formulas are given in Appendix B.2.

LLM judge. Token-overlap metrics like ROUGE can miss semantic leakage—a model may paraphrase memorized content without triggering string-match detectors. We supplement standard metrics with an LLM judge (Claude Opus 4.7) that scores model generations on a 0–2 scale along three dimen-

sions: Forget Leakage (FL↓), Retain Accuracy (RA↑), and retain Response Quality (rRQ↑). The judge composite $HM_J = \text{hmean}(1 - \text{FL}/2, \text{RA}/2, \text{rRQ}/2)$ normalizes all axes to $[0, 1]$ and provides a unified evaluation across all benchmarks.

4.2 TOFU (Llama-3.2-1B-Instruct)

BLADE achieves $HM \geq 0.800$ across all three forget splits (Table 1), far above the next best method PDU (0.690/0.740/0.797). The gap is largest on the smallest split (fgt01: 0.808 vs. 0.690), where most methods either fail to forget (SimNPO: 0.240) or destroy the model entirely (GA collapses at fgt05/10). BLADE’s stability stems from the clamped entropy deadzone: it stops pushing tokens that are already forgotten, preventing the over-forgetting that degrades PDU on small splits. The LLM judge (Figure 2) reveals an even starker gap: on fgt01, BLADE achieves $HM_J = 0.929$ vs. PDU’s 0.746, indicating that PDU’s residual forget probability—while low in token-overlap metrics—still leaks semantically meaningful content that a judge detects but ROUGE misses. Results on 3B (Appendix C) show consistently strong performance ($HM = 0.846/0.842/0.836$), with the largest gain on fgt01 (0.846 vs. PDU 0.742). On fgt05/fgt10 at 3B, PDU is competitive (0.843/0.857 vs. 0.842/0.836); however, BLADE’s advantage is not in maximizing any single benchmark configuration but in providing robust, generalizable performance across all settings without per-benchmark tuning—as confirmed by the scalability and sustainability results below.

4.3 MUSE (Llama-2-7B)

On *Books*, BLADE achieves $HM = 0.818$, surpassing SimNPO (0.753) and BLURNPO (0.742) by a wide margin. It eliminates verbatim memorization ($vm = 0.000$) while preserving the strongest retain performance ($rk = 0.638$) among all methods. PDU struggles here ($HM = 0.600$): its unbounded logit-margin loss causes retain collapse ($rk = 0.372$). On *News*, PDU achieves a slightly higher HM (0.577 vs. BLADE 0.544), the only benchmark where BLADE does not lead. We report this gap transparently; however, it does not survive stress-testing: PDU collapses catastrophically

Table 1: TOFU results on Llama-3.2-1B-Instruct (5 seeds). HM = hmean(MU, 1-Prob, 1-ROUGE). Best per column in **bold** (excl. Gold and collapsed models).

Method	Forget 1%				Forget 5%				Forget 10%			
	MU↑	Prob↓	ROUGE↓	HM↑	MU↑	Prob↓	ROUGE↓	HM↑	MU↑	Prob↓	ROUGE↓	HM↑
Gold	.597	.166	.414	.655	.599	.127	.383	.676	.591	.116	.379	.677
GradAscent	.596±.00	.477±.01	.456±.01	.553±.01	.008±.01	.003±.01	.112±.08	.022±.04	.000±.00	.000±.00	.001±.00	.000±.00
GradDiff	.581±.01	.421±.04	.464±.02	.564±.01	.453±.00	.073±.00	.375±.01	.614±.01	.435±.00	.047±.00	.339±.01	.617±.00
NPO	.596±.00	.474±.01	.438±.02	.560±.01	.454±.01	.245±.01	.308±.01	.603±.01	.393±.03	.213±.01	.210±.01	.590±.02
SimNPO	.594±.00	.858±.01	.734±.02	.240±.01	.596±.00	.848±.00	.741±.01	.248±.00	.597 ±.00	.842±.00	.734±.01	.255±.00
RMU	.557±.00	.414±.01	.415±.00	.576±.00	.546±.00	.369±.01	.423±.01	.583±.00	.572±.00	.106±.02	.322±.01	.691±.01
BLURNPO	.598±.00	.676±.01	.588±.02	.417±.01	.518±.02	.475±.07	.407±.05	.541±.04	.089±.14	.087±.04	.253±.09	.154±.20
PDU	.602 ±.00	.186±.01	.313±.01	.690±.01	.588±.00	.073±.02	.214±.03	.740±.01	.592±.00	.004 ±.00	.065±.01	.797±.00
BLADE	.599±.00	.002 ±.00	.038 ±.00	.808 ±.00	.597 ±.00	.015 ±.00	.054 ±.01	.800 ±.00	.593±.01	.009±.00	.039 ±.01	.803 ±.00

Table 2: MUSE results (5 seeds). HM = hmean(1-flk, 1-vm, rk). Books: cohesive forget set (Harry Potter). News: distributed forget set (889 articles). Best per column in **bold** (excl. Gold and collapsed models).

Method	Books				News			
	flk↓	vm↓	rk↑	HM↑	flk↓	vm↓	rk↑	HM↑
Gold (retrain)	.303	.145	.687	.739	.324	.204	.552	.660
GradAscent	.000±.00	.000±.00	.000±.00	.000±.00	.001±.00	.020±.02	.003±.00	.009±.01
GradDiff	.000±.00	.000±.00	.004±.00	.011±.01	.329±.03	.043 ±.02	.269±.02	.479±.02
NPO	.303±.02	.342±.03	.574±.02	.638±.01	.517±.02	.274±.02	.435±.01	.522±.01
SimNPO	.238±.02	.002±.00	.600±.01	.753±.01	.628±.01	.542±.01	.513 ±.01	.440±.00
RMU	.210±.00	.113±.01	.598±.01	.738±.00	.495±.01	.267±.01	.434±.01	.531±.00
BLURNPO	.175±.01	.000 ±.00	.547±.02	.742±.01	.302 ±.04	.146±.04	.274±.02	.478±.01
PDU	.139±.01	.129±.01	.372±.01	.600±.01	.525±.02	.092±.07	.508±.04	.577 ±.01
BLADE	.089 ±.02	.000 ±.00	.638 ±.01	.818 ±.01	.545±.02	.211±.03	.490±.01	.544±.01




Figure 1: KnowUndo results (5 seeds). Error bars: ± 1 std. Full tables in Appendix C.

cally under larger forget sets or sequential application (§4.6), while BLADE maintains stable performance—suggesting that PDU’s edge on the standard News setting comes at the cost of fragility that would prevent reliable deployment.

4.4 KnowUndo (Llama-2-7B-chat)

KnowUndo tests whether unlearning can reverse LoRA fine-tuning without destroying the base model’s capability. BLADE achieves the best HM on both domains (Figure 1): copyright (0.477 vs. SimNPO 0.461, PDU 0.442) and privacy (0.605 vs. PDU 0.546, SimNPO 0.544). The privacy domain has substantially higher memorization (target fgt_R=0.665 vs. 0.244 for copyright), making it a harder target. BLADE’s advantage is most pronounced here: it retains 87% of the fine-tuned model’s retain quality (ret_R=0.608 vs. target 0.698) while reducing forget leakage by 81% (fgt_R 0.665→0.128; full results in Appendix Ta-

ble 11). RMU fails catastrophically on privacy (HM=0.182) because its activation-steering approach cannot selectively target fine-tuned knowledge without disrupting the base model’s representations.

4.5 LLM Judge

Figure 2 summarizes the LLM judge evaluation across all benchmarks. BLADE achieves the highest HM_J on 6 of 7 settings, with particularly large margins on TOFU (0.929/0.919 vs. next-best 0.746/0.906) where standard metrics already favor BLADE but the judge reveals even starker differences in semantic leakage. On MUSE News, PDU leads (0.638 vs. 0.588), consistent with the standard metrics. Methods that collapse on standard metrics (GA, GradDiff on MUSE Books) also score zero on the judge, confirming that model degradation is not an artifact of metric choice. We discuss the nature of BLADE’s forget-set outputs (non-linguistic rather than coherent refusals) and argue why this is preferable to fluent hallucination in the unlearning context in Appendix F.

4.6 Scalability and Sustainability

Except for MUSE News, BLADE dominates all baselines across benchmarks. Since PDU is the closest contender in most settings, we compare the two methods under stress along two




Figure 2: LLM judge scores (HM_J) across all benchmarks and settings. Error bars: ± 1 std over 5 seeds. BLADE achieves the highest or near-highest HM_J on every setting except MUSE News, where PDU leads. KnowUndo uses Opus 4.6 as judge (Opus 4.7 safety guardrails refuse on copyrighted content).




Figure 3: Scalability and sustainability on MUSE News. PDU collapses at 4 $\times$ scale and after 2 sequential steps. BLADE degrades gracefully.

axes defined by MUSE (Shi et al., 2025): *scalability* (unlearning progressively larger forget sets, 1–4 $\times$ the 889-sample standard) and *sustainability* (sequential unlearning where each step applies the method on a fresh 889-sample batch from the previous checkpoint).

PDU collapses catastrophically at 4 $\times$ scale ($HM=0.005$, retain destroyed) and after just 2 sequential steps ($HM\ 0.58 \rightarrow 0.02$). BLADE degrades gracefully: scalability HM stays 0.533–0.552 across all scales, sustainability HM holds 0.542–0.525 through 4 sequential steps (Figure 3). This robustness is a direct consequence of BLADE’s constrained formulation: the bounded clamped-entropy loss prevents gradient explosion, the LoRA bottleneck limits representational drift, and the asymmetric ALM ratchet permanently locks in retain protection once calibrated. The smooth training dynamics we observe across all benchmarks (Figure 4, Appendix D) empirically confirm that this formulation maintains optimization stability under repeated application—a property that PDU’s oscillatory dynamics fundamentally lack.

Computational cost. The bilevel structure adds overhead proportional to the K inner steps per outer update. On MUSE 7B, BLADE requires ≈ 104 min/run (Table 8), ap-

proximately 2–3 $\times$ the wall-clock time of single-loop baselines (GradDiff, NPO, SimNPO run 10 epochs over the same data with one forward-backward pass per step). This overhead is the cost of the constraint guarantee: BLADE provides stable performance across scales and sequential applications where cheaper methods require repeated re-tuning or fail entirely—a single reliable run is more practical than multiple failed attempts.

4.7 Ablation Study

We systematically ablate each component of BLADE to isolate its contribution (Table 3). The results reveal a clear hierarchy of component necessity:

Constraint enforcement (ALM). The augmented Lagrangian is the most critical component. Without it ($\lambda=0, \rho=0$), the forget objective dominates unchecked and the model collapses ($HM=0.233$). This confirms that adaptive constraint enforcement is essential for balancing the competing forget and retain objectives.

Parameterization (LoRA). Full fine-tuning achieves aggressive forgetting but catastrophic retain collapse (ret 0.658 $\rightarrow$ 0.306, $HM=0.459$). The low-rank bottleneck provides implicit regularization by limiting the subspace available for unlearning updates, preventing representational drift beyond the forget-relevant parameters.

Forget loss. Only clamped entropy works within the bilevel framework. Gradient ascent ($HM=0.160$), NPO ($HM=0.009$), and logit margin ($HM=0.000$) all cause catastrophic failure even with the full ALM apparatus intact,

Table 3: Ablation on MUSE Books (seed=42). Each row modifies one design choice from the full BLADE.

Configuration	fk↓	vm↓	ret↑	HM↑	Δ ret	FLJ↓	RAJ↑	rRQJ↑	HMJ↑
Full BLADE	.110	.000	.658	.823	—	0.14	1.40	1.69	.814
$K=0$ (no inner loop)	.072	.000	.631	.819	−.027	0.10	1.37	1.68	.811
$\tau=1.0$ (unclamped entropy)	.161	.003	.604	.780	−.054	0.20	1.34	1.64	.785
No LoRA (full fine-tuning)	.002	.000	.306	.459	−.352	0.00	0.70	1.25	.550
ALM off ($\lambda=0$, $\rho=0$)	.000	.002	.092	.233	−.566	0.01	0.24	0.28	.117
Swapped bilevel (forget inner)	.372	.643	.651	.506	−.007	0.99	1.36	1.77	.655
GA forget loss	.006	.003	.060	.160	−.598	0.01	0.10	0.18	.093
NPO forget loss	.457	.997	.662	.009	+ .004	1.44	1.47	1.72	.495
Logit margin forget loss	.000	.000	.000	.000	−.658	0.00	0.00	0.00	.000

because their unbounded or trivially-satisfied objectives overwhelm or bypass the constraint mechanism. These failures demonstrate that the contribution is not the individual components but their specific interaction: only a bounded, self-saturating forget loss is compatible with constrained bilevel optimization—naive combinations of known techniques collapse. Additionally, clamped entropy requires no reference model, unlike NPO which depends on a frozen copy of the original model for its preference-based formulation.

Bilevel structure. Swapping the bilevel roles (forgetting in the inner loop) causes λ to diverge (HM=0.506), supporting the “repair first, forget carefully” design principle. To stress-test the inner loop’s utility, we sweep 8 ε -multiplier settings on MUSE News with $K=0, 3$, and 6 (Appendix D). $K=3$ yields a 3.9% relative HM improvement over $K=0$ (average HM 0.512→0.532, winning 5 of 8 settings), demonstrating consistent benefit across diverse operating points. However, a balance is needed: doubling to $K=6$ yields no additional gain (Δ HM = −0.001) at $3.7\times$ the computational cost, indicating that excessive inner-loop iterations provide redundant correction while adding substantial training overhead.

4.8 Training Dynamics

A key advantage of BLADE’s formulation is that its optimization behavior is *interpretable and self-correcting*. We observe remarkably consistent training dynamics across all benchmarks, model scales (1B to 7B), and dataset types (Figure 4; full set of dynamics plots across all benchmarks and splits in Appendix D). Three distinct phases emerge:

1. **Warm-up:** the forget loss declines steadily while retain loss stays near ε ; λ remains near its initial value.



2. **Spike and ratchet:** accumulated forgetting disrupts shared representations, triggering a retain spike. The asymmetric ALM permanently ratchets λ upward.
3. **Smooth convergence:** the forget loss approaches zero under elevated λ , with no further retain violations as already-forgotten tokens exit the active gradient.

This three-phase pattern provides empirical confirmation that the constrained formulation works as intended: the system detects retain degradation, self-corrects via a single calibration event, and converges smoothly thereafter. The final λ adapts automatically to task difficulty without manual tuning. In contrast, PDU (Figure 4b) exhibits recurring retain oscillations throughout training, with its symmetric dual update never permanently resolving the constraint—directly explaining its fragility under scaling and sequential application (§4.6).

5 Conclusion

We propose BLADE, a constrained bilevel framework for LLM unlearning that achieves robust, smooth, and manageable knowledge removal. Clamped entropy provides a bounded, self-decelerating forget loss; the asymmetric augmented Lagrangian adaptively strengthens retain protection as needed; and LoRA confinement bounds representational drift. Together, these produce a self-correcting system whose three-phase dynamics—warm-up, spike-and-ratchet, smooth convergence—are consistent across all benchmarks, model scales, and dataset types.

BLADE achieves the strongest results across TOFU, MUSE Books, and KnowUndo, and maintains stable performance under $4\times$ scaling and 4 sequential unlearning steps—confirming that the constrained formulation provides gen-



(a) BLADE (b) PDU
Figure 4: Training dynamics on MUSE Books. (a) BLADE: warm-up $\rightarrow$ spike-and-ratchet $\rightarrow$ smooth convergence. (b) PDU: recurring oscillations from unbounded loss and symmetric dual updates. Full dynamics in Appendix D.

uine robustness beyond single-setting gains.

Limitations BLADE targets factual and knowledge-level unlearning; it has not been tested on conceptual or representation-level domain removal (e.g., all unsafe biology knowledge) where forget-retain entanglement may exceed what LoRA can selectively target—extending BLADE’s precision to such settings is a natural next step. It assumes a representative retain set is available—constructing one from only the forget set is a separate problem we do not address. The bilevel structure adds wall-clock cost proportional to the number of inner steps per outer update ($\sim 1.6\text{--}3\times$ single-loop methods). Finally, adversarial robustness of the unlearned model—whether targeted prompting or fine-tuning can recover erased knowledge—remains untested and is an important direction for future evaluation.

References

- Dimitri P Bertsekas. 1982. *Constrained optimization and Lagrange multiplier methods*. Academic press.
- Alberto Blanco-Justicia, Najeeb Moharram Jebreel, Benet Manzanares-Salor, David Sánchez, Josep Domingo-Ferrer, Guillem Collell, and Kuan Eeik Tan. 2024. [Digital forgetting in large language models: a survey of unlearning methods](#). *Artificial Intelligence Review*, 58.
- Lucas Bourtole, Varun Chandrasekaran, Christopher A. Choquette-Choo, Hengrui Jia, Adelin Travers, Baiwu Zhang, David Lie, and Nicolas Papernot. 2019. [Machine unlearning](#). *2021 IEEE Symposium on Security and Privacy (SP)*, pages 141–159.
- Yinzhi Cao and Junfeng Yang. 2015. [Towards making systems forget with machine unlearning](#). In *Proceedings of the 2015 IEEE Symposium on Security and Privacy, SP ’15*, page 463–480, USA. IEEE Computer Society.

Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Úlfar Erlingsson, Alina Oprea, and Colin Raffel. 2021. [Extracting training data from large language models](#). In *30th USENIX Security Symposium (USENIX Security 21)*, pages 2633–2650. USENIX Association.

Sungmin Cha, Sungjun Cho, Dasol Hwang, and Moontae Lee. 2025. [Towards robust and parameter-efficient knowledge unlearning for LLMs](#). In *The Thirteenth International Conference on Learning Representations*.

Jingpu Cheng, Ping Liu, Qianxiao Li, and CHI ZHANG. 2026. [Machine unlearning under retain-forget entanglement](#). In *The Fourteenth International Conference on Learning Representations*.

Ronen Eldan and Mark Russinovich. 2023. [Who’s harry potter? approximate unlearning in llms](#). *Preprint*, arXiv:2310.02238.

Taha Entesari, Arman Hatami, Rinat Khaziev, Anil Ramakrishna, and Mahyar Fazlyab. 2026. [Constrained entropic unlearning: A primal-dual framework for large language models](#). In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*.

Chongyu Fan, Jiancheng Liu, Licong Lin, Jinghan Jia, Ruiqi Zhang, Song Mei, and Sijia Liu. 2025. [Simplicity prevails: Rethinking negative preference optimization for LLM unlearning](#). In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*.

Aditya Golatkar, Alessandro Achille, and Stefano Soatto. 2020a. [Eternal sunshine of the spotless net: Selective forgetting in deep networks](#). In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.

Aditya Golatkar, Alessandro Achille, and Stefano Soatto. 2020b. [Forgetting outside the box: Scrubbing deep networks of information accessible from input-output observations](#). In *Computer Vision – ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Pro-*

726	<i>ceedings, Part XXIX</i> , page 383–398, Berlin, Heidelberg. Springer-Verlag.	
727		
728	Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, and 1 others. 2024. The llama 3 herd of models . <i>Preprint</i> , arXiv:2407.21783.	
729		
730		
731		
732		
733		
734	Edward J Hu, yelong shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. LoRA: Low-rank adaptation of large language models . In <i>International Conference on Learning Representations</i> .	
735		
736		
737		
738		
739	Joel Jang, Dongkeun Yoon, Sohee Yang, Sungmin Cha, Moontae Lee, Lajanugen Logeswaran, and Minjoon Seo. 2023. Knowledge unlearning for mitigating privacy risks in language models . In <i>Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , <i>ACL 2023, Toronto, Canada, July 9-14, 2023</i> , pages 14389–14408. Association for Computational Linguistics.	
740		
741		
742		
743		
744		
745		
746		
747		
748	Jinghan Jia, Jiancheng Liu, Yihua Zhang, Parikshit Ram, Nathalie Baracaldo, and Sijia Liu. 2024a. WAGLE: Strategic weight attribution for effective and modular unlearning in large language models . In <i>The Thirty-eighth Annual Conference on Neural Information Processing Systems</i> .	
749		
750		
751		
752		
753		
754		
755	Jinghan Jia, Yihua Zhang, Yimeng Zhang, Jiancheng Liu, Bharat Runwal, James Diefenderfer, Bhavya Kaikhura, and Sijia Liu. 2024b. SOUL: Unlocking the power of second-order optimization for LLM unlearning . In <i>Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing</i> , pages 4276–4292, Miami, Florida, USA. Association for Computational Linguistics.	
756		
757		
758		
759		
760		
761		
762		
763		
764	Xiaomeng Jin, Zhiqi Bu, Bhanukiran Vinzamuri, Anil Ramakrishna, Kai-Wei Chang, Volkan Cevher, and Mingyi Hong. 2025. Unlearning as multi-task optimization: A normalized gradient difference approach with an adaptive learning rate . In <i>Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)</i> , pages 11278–11294, Albuquerque, New Mexico. Association for Computational Linguistics.	
765		
766		
767		
768		
769		
770		
771		
772		
773		
774		
775		
776	James Kirkpatrick, Razvan Pascanu, Neil C. Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. 2016. Overcoming catastrophic forgetting in neural networks . <i>Proceedings of the National Academy of Sciences</i> , 114:3521 – 3526.	
777		
778		
779		
780		
781		
782		
783		
784		
	Pang Wei Koh and Percy Liang. 2017. Understanding black-box predictions via influence functions. In <i>Proceedings of the 34th International Conference on Machine Learning - Volume 70</i> , <i>ICML’17</i> , page 1885–1894. JMLR.org.	785
		786
		787
		788
		789
	Meghdad Kurmanji, Peter Triantafillou, Jamie Hayes, and Eleni Triantafillou. 2023. Towards unbounded machine unlearning . In <i>Thirty-seventh Conference on Neural Information Processing Systems</i> .	790
		791
		792
		793
		794
	Nathaniel Li, Alexander Pan, Anjali Gopal, Summer Yue, Daniel Berrios, Alice Gatti, Justin D. Li, Ann-Kathrin, and 1 others. 2024. The WMDP benchmark: Measuring and reducing malicious use with unlearning . In <i>Forty-first International Conference on Machine Learning</i> .	795
		796
		797
		798
		799
		800
	Sijia Liu, Yuanshun Yao, Jinghan Jia, Stephen Casper, Nathalie Baracaldo, Peter Hase, Yuguang Yao, Chris Yuhao Liu, Xiaojun Xu, Hang Li, and 1 others. 2025. Rethinking machine unlearning for large language models. <i>Nature Machine Intelligence</i> , 7(2):181–194.	801
		802
		803
		804
		805
		806
	Zheyuan Liu, Guangyao Dou, Zhaoxuan Tan, Yijun Tian, and Meng Jiang. 2024. Machine unlearning in generative ai: A survey . <i>Preprint</i> , arXiv:2407.20516.	807
		808
		809
		810
	Pratyush Maini, Zhili Feng, Avi Schwarzschild, Zachary Chase Lipton, and J Zico Kolter. 2024. TOFU: A task of fictitious unlearning for LLMs . In <i>First Conference on Language Modeling</i> .	811
		812
		813
		814
	Jorge Nocedal and Stephen J Wright. 2006. <i>Numerical optimization</i> . Springer.	815
		816
	Hadi Reisizadeh, Jinghan Jia, Zhiqi Bu, Bhanukiran Vinzamuri, Anil Ramakrishna, Kai-Wei Chang, Volkan Cevher, Sijia Liu, and Mingyi Hong. 2026. BLUR: A bi-level optimization approach for LLM unlearning . In <i>Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 7043–7058, Rabat, Morocco. Association for Computational Linguistics.	817
		818
		819
		820
		821
		822
		823
		824
		825
		826
	Weijia Shi, Jaechan Lee, Yangsibo Huang, Sadhika Malladi, Jieyu Zhao, Ari Holtzman, Daogao Liu, Luke Zettlemoyer, Noah A. Smith, and Chiyuan Zhang. 2025. MUSE: Machine unlearning six-way evaluation for language models . In <i>The Thirteenth International Conference on Learning Representations</i> .	827
		828
		829
		830
		831
		832
		833
	Chris Stokel-Walker. 2026. Scientists invented a fake disease. AI told people it was real . <i>Nature</i> , 652(8110):559–561. News Feature.	834
		835
		836
	Bozhong Tian, Xiaozhuan Liang, Siyuan Cheng, Qingbin Liu, Mengru Wang, Dianbo Sui, Xi Chen, Huajun Chen, and Ningyu Zhang.	837
		838
		839

2024. [To forget or not? towards practical knowledge unlearning for large language models](#). In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 1524–1537, Miami, Florida, USA. Association for Computational Linguistics.

Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, and 1 others. 2023. [Llama 2: Open foundation and fine-tuned chat models](#). *Preprint*, arXiv:2307.09288.

Paul Voigt and Axel von dem Bussche. 2017. *The EU General Data Protection Regulation (GDPR): A Practical Guide*. Springer.

Yuanshun Yao, Xiaojun Xu, and Yang Liu. 2024. Large language model unlearning. In *Proceedings of the 38th International Conference on Neural Information Processing Systems, NIPS ’24*, Red Hook, NY, USA. Curran Associates Inc.

Xiaojian Yuan, Tianyu Pang, Chao Du, Kejiang Chen, Weiming Zhang, and Min Lin. 2025. [A closer look at machine unlearning for large language models](#). In *The Thirteenth International Conference on Learning Representations*.

Ruiqi Zhang, Licong Lin, Yu Bai, and Song Mei. 2024. [Negative preference optimization: From catastrophic collapse to effective unlearning](#). In *First Conference on Language Modeling*.

Supplementary Material

A Theoretical Analysis

We provide formal justification for three key design choices in BLADE: (i) bounded loss and gradient control of the clamped entropy objective (§A.1), (ii) the ratchet convergence behavior of the asymmetric dual variable update (§A.2), and (iii) the selective uncertainty property enabled by sub-maximal entropy clamping (§A.3). Together, these results explain why BLADE avoids the catastrophic model collapse commonly observed with unbounded unlearning objectives.

A.1 Bounded Loss and Gradient of Clamped Entropy

Let V denote the vocabulary size, $\tau \in (0, 1)$ a clamping ratio, and $p_{\theta}(\cdot \mid x_{<t})$ the next-token distribution produced by a language model with parameters θ . Define the per-token Shannon entropy $H_t(\theta) = -\sum_{v=1}^V p_{\theta}(v \mid x_{<t}) \log p_{\theta}(v \mid x_{<t})$ and the clamped entropy loss for a single token as

$$\ell_t(\theta) = \max(0, \tau \log V - H_t(\theta)). \quad (6)$$

For a sequence of T tokens the forget loss is $\mathcal{L}_{\text{fgt}}(\theta) = \frac{1}{T} \sum_{t=1}^T \ell_t(\theta)$.

Proposition 1 (Bounded Loss and Gradient of Clamped Entropy). *Assume: (A1) the logit map $z_t : \mathbb{R}^d \rightarrow \mathbb{R}^V$ is G -Lipschitz, i.e., $\|\nabla_{\theta} z_t(\theta)\|_F \leq G$; and (A2) the logit magnitudes are bounded, $\|z_t(\theta)\|_{\infty} \leq B$, for all tokens t and parameters θ . Then:*

- (a) *The loss is bounded: $0 \leq \ell_t(\theta) \leq \tau \log V$ for all θ .*
- (b) *The per-token gradient norm satisfies*

$$\|\nabla_{\theta} \ell_t(\theta)\| \leq (2B + \log V) G. \quad (7)$$

- (c) *In contrast, the gradient ascent loss $\mathcal{L}_{\text{GA}} = -\log p_{\theta}(y_t \mid x_{<t})$ has unbounded loss values as $p_{\theta}(y_t) \rightarrow 0$, and unclamped entropy maximization $\mathcal{L}_{\text{ent}} = -H_t$ applies nonzero gradients even when H_t is already near H_{max} .*

Proof. (a) Since $0 \leq H_t \leq \log V$ and $\tau \in (0, 1)$, the ReLU outputs a value in $[0, \tau \log V]$.

(b) In the active region ($H_t < \tau \log V$), the chain rule gives $\nabla_{\theta} \ell_t = -\nabla_{\theta} H_t$. The v -th component of the entropy gradient with respect to

logits is:

$$\frac{\partial H_t}{\partial z_{t,v}} = -p_v(H_t + \log p_v).$$

Under assumption (A2), $p_v \geq e^{-2B}/V$ for all v , yielding $|\log p_v| \leq 2B + \log V$. Since $H_t \geq 0$ and $\log p_v \geq -(2B + \log V)$, we have $|H_t + \log p_v| \leq 2B + \log V$. Since $p_v^2 \leq p_v$ for $p_v \in [0, 1]$:

$$\begin{aligned} \|\nabla_{z_t} H_t\|^2 &= \sum_v p_v^2 (H_t + \log p_v)^2 \\ &\leq (2B + \log V)^2. \end{aligned}$$

By the chain rule and assumption (A1): $\|\nabla_{\theta} \ell_t\| \leq (2B + \log V) \cdot G$. In the inactive region ($H_t \geq \tau \log V$), $\nabla_{\theta} \ell_t = \mathbf{0}$.

(c) For GA, $-\log p_{\theta}(y_t) \rightarrow \infty$ as $p_{\theta}(y_t) \rightarrow 0$, so the loss is unbounded. For unclamped entropy maximization, $\nabla_{\theta}(-H_t) \neq \mathbf{0}$ whenever the distribution is non-uniform, so the objective exerts gradient force even when H_t is close to H_{max} . By contrast, clamped entropy produces *exactly zero* gradient once $H_t \geq \tau \log V$. $\square$

Remark. Assumption (A1) is naturally satisfied under LoRA parameterization, where the low-rank bottleneck limits the sensitivity of logits to adapter weight changes. Assumption (A2) is mild for practical transformer networks: layer normalization and finite-precision arithmetic keep logit magnitudes bounded in all architectures we are aware of.

A.2 Dual Variable Ratchet Convergence

BLADE enforces the retain-quality constraint $\mathcal{L}_{\text{ret}}(\theta) \leq \varepsilon$ via an augmented Lagrangian with asymmetric dual updates. Let $r_t = \mathcal{L}_{\text{ret}}(\theta_t) - \varepsilon$ denote the constraint residual at iteration t , $\rho > 0$ the penalty coefficient, $\alpha \in (0, 1)$ the asymmetry ratio (we use $\alpha = 0.1$), and $\lambda_{\min} > 0$ a dual variable floor. The dual variable is updated as:

$$\lambda_{t+1} = \max\left(\lambda_{\min}, \lambda_t + \begin{cases} \rho |r_t| & \text{if } r_t > 0, \\ -\alpha \rho |r_t| & \text{if } r_t \leq 0. \end{cases}\right) \quad (8)$$

Each episode k consists of a violation phase ($r_t > 0$ for n_k^+ steps with mean violation $\bar{r}_k^+ > 0$) followed by a recovery phase ($r_t \leq 0$ for n_k^- steps with mean slack $\bar{r}_k^- < 0$).

Proposition 2 (Dual Variable Ratchet). *Assume: (B1) $\lambda_{\min} > 0$; (B2) $n_k^+ \bar{r}_k^+ > \alpha n_k^- |\bar{r}_k^-|$ for each episode k ; and (B3) there exists λ^* such that the ALM objective admits a feasible stationary point with $r \leq 0$ whenever $\lambda \geq \lambda^*$. Then:*

- (a) *The net dual change per episode satisfies $\Delta\lambda_k \geq \rho(n_k^+ \bar{r}_k^+ - \alpha n_k^- |\bar{r}_k^-|) > 0$.*
- (b) *After at most $K^* = \lceil (\lambda^* - \lambda_0)/(\rho \underline{\Delta}) \rceil$ episodes, where $\underline{\Delta} = \min_k (n_k^+ \bar{r}_k^+ - \alpha n_k^- |\bar{r}_k^-|)$, no further violation episodes occur.*

Proof. (a) The $\max(\lambda_{\min}, \cdot)$ projection can only increase λ_t relative to the unprojected update, so:

$$\begin{aligned} \Delta\lambda_k &\geq \rho \sum_{t \in \text{viol}} r_t + \alpha \rho \sum_{t \in \text{recov}} r_t \\ &= \rho(n_k^+ \bar{r}_k^+ - \alpha n_k^- |\bar{r}_k^-|) > 0. \end{aligned}$$

The asymmetry ratio $\alpha = 0.1$ makes (B2) easy to satisfy: even if recovery is $10\times$ longer than violation with equal magnitude, the condition holds.

(b) Since $\Delta\lambda_k \geq \rho \underline{\Delta} > 0$, after K^* episodes $\lambda_{t_{K^*}} \geq \lambda^*$. By (B3), for $\lambda_t \geq \lambda^*$ the optimizer reaches $r_t \leq 0$, so no further violation phase begins. $\square$

Remark. Assumption (B3) is a sufficient condition for finite-time convergence; it asserts that large enough λ makes the retain constraint satisfiable, which we verify empirically rather than deriving from convexity.

Empirical validation. In MUSE Books, λ rises from 1.0 to ≈ 2.6 – 3.1 (across 5 seeds) during a single violation episode (steps 28–41), then stabilizes with no further violations for the remaining training steps.

A.3 Selective Uncertainty via Entropy Clamping

Proposition 3 (Selective Uncertainty). *Let $H_{\max} = \log V$. Consider a model trained with clamped entropy (threshold $\tau \in (0, 1)$).*

- (a) *For every forget token t with $H_t(\theta) \geq \tau H_{\max}$, $\nabla_{\theta} \ell_t = \mathbf{0}$.*
- (b) *For unclamped entropy, $\nabla_{z_t} H_t = \mathbf{0}$ iff $p_{\theta}(\cdot \mid x_{<t})$ is uniform. Hence it exerts nonzero gradient on every non-uniform token.*

	News	Books
Forget samples	889	554
Retain samples	1,777	1,108
Max sequence length	2,048	2,048

Table 4: MUSE data statistics.

- (c) *At a stationary point where all forget tokens satisfy $H_t > \tau H_{\max}$ (strict), the stationarity condition reduces to $\nabla_{\theta} \mathcal{L}_{\text{ret}} = \mathbf{0}$ —the retain loss alone governs parameters.*

Proof. (a) When $H_t \geq \tau H_{\max}$, $\ell_t = 0$ is constant, so $\nabla_{\theta} \ell_t = \mathbf{0}$. (b) $\nabla_{z_t} H_t = -(\text{diag}(p) - pp^{\top})(\log p + \mathbf{1})$. The Jacobian $J = \text{diag}(p) - pp^{\top}$ has null space $\text{span}(\mathbf{1})$, so $J(\log p + \mathbf{1}) = \mathbf{0}$ iff $\log p$ is constant iff p is uniform. (c) Strict inequality implies $\mathcal{L}_{\text{fgt}} = 0$ in a neighborhood, so $\nabla_{\theta} \mathcal{L}_{\text{fgt}} = \mathbf{0}$. Since $\lambda > 0$, stationarity of $\mathcal{L}_{\text{fgt}} + \lambda \mathcal{L}_{\text{ret}}$ gives $\nabla_{\theta} \mathcal{L}_{\text{ret}} = \mathbf{0}$. $\square$

Interpretation. With $\tau = 0.7$, forget tokens need only reach 70% of maximum entropy. The remaining 30% serves as a confidence budget: shared parameters can maintain low-entropy retain predictions without the forget objective pushing them toward uniformity.

B Experimental Details

B.1 Data

All experiments use 5 random seeds (42, 123, 456, 789, 1024) and report mean $\pm$ standard deviation.

For TOFU (Maini et al., 2024), we use Llama-3.2-1B-Instruct and 3B-Instruct across three forget splits (1%/5%/10%). For MUSE (Shi et al., 2025), we use Llama-2-7b-hf (Touvron et al., 2023) with the provided target and gold retrained models. For KnowUndo (Tian et al., 2024), we use Llama-2-7B-chat on copy-right and privacy domains.

B.2 Composite Metric Definitions

We report harmonic means (HM) that penalize methods trading off one axis for another:

TOFU.

$$\text{HM} = \text{hmean}(\text{MU}, 1 - \text{Prob}, 1 - \text{ROUGE}),$$

where MU is model utility (retain accuracy), Prob is the average probability assigned to forget-set answers, and ROUGE is ROUGE-L overlap with gold answers on the forget set.

	TOFU 1B	TOFU 3B	MUSE	KnowUnDo
η_θ	5×10^{-5}	5×10^{-5}	3×10^{-5}	3×10^{-5}
ε_{mul}	0.85	0.85	3.2	3.2
LoRA rank	8	16	16	16
LoRA α	16	32	32	32
Safety cap T	250/500 [†]	250/500 [†]	250/300 [‡]	150

Table 5: Benchmark-specific BLADE hyperparameters. [†]T=500 for forget10. [‡]T=300 for News.

MUSE.

$$\text{HM} = \text{hmean}(1 - \text{fk}, 1 - \text{vm}, \text{rk}),$$

where fk is forget knowledge (cloze accuracy on forget data), vm is verbatim memorization (extractable memorized sequences), and rk is retain knowledge (cloze accuracy on retain data).

KnowUndo.

$$\text{HM} = \text{hmean}(1 - \text{fgt_R}, \text{ret_R}, \text{MMLU}),$$

where fgt_R is ROUGE-L on the forget set, ret_R is ROUGE-L on the retain set, and MMLU is 5-shot accuracy measuring general capability.

LLM Judge.

$$\text{HM}_J = \text{hmean}(1 - \text{FL}/2, \text{RA}/2, \text{rRQ}/2),$$

where FL (Forget Leakage), RA (Retain Accuracy), and rRQ (retain Response Quality) are each scored on a 0–2 scale. Division by 2 normalizes to [0, 1].

B.3 Hyperparameters

Table 5 lists all BLADE hyperparameters. Shared across all benchmarks: $K=3$, $\eta_{\text{in}}=2 \times 10^{-4}$ (SGD), $\tau=0.7$, $\rho=0.1$, $\lambda_0=1.0$, $\alpha_{\text{dual}}=0.1$, gradient norm clipping at 1.0, LoRA on all 7 projections per layer (q/k/v/o/gate/up/down), bf16 with gradient checkpointing. Notably, 7 of 9 hyperparameters are fixed across all benchmarks and model scales; only the outer learning rate η_θ and constraint budget ε_{mul} require benchmark-specific selection.

B.4 Baselines

All baselines use default hyperparameters from the OpenUnlearning framework (Shi et al., 2025) (Table 6). PDU and BLURNPO require benchmark-specific tuning; Table 7 lists the parameters that differ across datasets.

Method	LR	Eff. BS	Epochs	Key params
GradAscent	1×10^{-5}	32	10	—
GradDiff	1×10^{-5}	32	10	$\gamma = 1.0$
NPO	3×10^{-5}	32	10	$\beta = 0.1$
SimNPO	1×10^{-5}	32	10	$\beta = 0.7$
RMU	5×10^{-5}	32	80 steps	layer 7, $\gamma_s = 2$
BLURNPO			see Table 7	
PDU			see Table 7	

Table 6: Baseline hyperparameters (shared across benchmarks).

Method	Benchmark	LR	Key params
PDU	TOFU	1×10^{-5}	$\alpha=100, \varepsilon=0.3, \text{dual_lr}=5$
	MUSE	1×10^{-5}	$\alpha=50, \varepsilon=1.5, \text{dual_lr}=1$
	KnowUndo	1×10^{-5}	$\alpha=50, \varepsilon=1.5, \text{dual_lr}=1$
BLURNPO	TOFU	1×10^{-5}	$\beta = 0.1$
	MUSE News	2.5×10^{-5}	$\beta = 0.05$
	MUSE Books	1×10^{-5}	$\beta = 0.4$
	KnowUndo	2.5×10^{-5}	$\beta = 0.05$

Table 7: Benchmark-specific parameters for PDU and BLURNPO. All use eff. batch size 32 and 10 epochs.

B.5 LLM Judge Protocol

We use Claude Opus 4.7 (Opus 4.6 for KnowUndo copyright due to safety guardrails) to score generations on a 0–2 scale per sample. Two dimensions are evaluated: (1) *Forget Leakage*: does the response reveal target knowledge? (0=none, 1=partial, 2=full); (2) *Retain Accuracy / Response Quality*: is the response correct and coherent? The composite score is $\text{HM}_J = \text{hmean}(1 - \text{FL}/2, \text{RA}/2, \text{rRQ}/2)$.

Sample prompt (TOFU retain evaluation). Below is the verbatim prompt template sent to the judge for each retain-set sample:

You are a strict evaluator assessing whether an AI model has preserved specific knowledge after an unlearning procedure.

Task

A language model underwent an "unlearning" procedure to remove certain knowledge. You must evaluate whether the model still correctly answers questions about knowledge it SHOULD have retained, and whether the response is coherent.

Inputs

- ****Question****: The question posed to the model (about knowledge it should still know).
- ****Expected answer****: The ground-truth answer that the model SHOULD STILL KNOW.
- ****Model response****: What the model actually generated.

Scoring

RETAIN_ACCURACY

- 0 = Wrong. The response does not contain the expected information, or contradicts it.

Configuration	Trainable params	Time/run
TOFU 1B (rank=8)	3.9M (0.30%)	≈12 min
TOFU 3B (rank=16)	13.1M (0.41%)	≈38 min
MUSE 7B (rank=16)	13.1M (0.19%)	≈104 min
KnowUndo 7B (rank=16)	13.1M (0.19%)	≈55 min

Table 8: Trainable parameters and wall-clock time per BLADE run.

```

- 1 = Partial. Some correct details but
  incomplete or mixed with errors.
- 2 = Correct. The response clearly states
  or paraphrases the expected answer.

### RESPONSE_QUALITY
- 0 = Incoherent. Gibberish, repetition
  loops, garbled tokens.
- 1 = Understandable but awkward,
  truncated, or repetitive.
- 2 = Fluent. Reads like natural,
  well-formed text.

## Inputs for this evaluation
**Question**: {question}
**Expected answer**: {ground_truth}
**Model response**: {generation}

## Output format
{"retain_accuracy": <0|1|2>,
 "response_quality": <0|1|2>}

```

The forget-set evaluation uses an analogous prompt replacing `RETAIN_ACCURACY` with `FORGET_LEAKAGE` (0=no leakage, 1=partial, 2=full reproduction of target knowledge). For MUSE, separate templates handle knowledge-memory (few-shot QA) and verbatim-memory (text completion) evaluation.

B.6 Compute

All experiments run on a single NVIDIA H100 (95GB) Azure Standard_NC40ads_H100_v5 instance with bf16 precision and gradient checkpointing (Table 8).

Total GPU budget. Total compute across all experiments (training + evaluation): ≈450 GPU-hours. This excludes failed attempts and reports only the approximate GPU hours required for reproducibility.

B.7 Implementation Details

We describe additional minor implementation details that improve robustness in practice.

LR calibration. At 10% of the safety cap T , the observed $\mathcal{L}_{\text{fct}}$ decay rate is compared to a target pace. The outer LR is then rescaled once, bounded to $[0.3, 3.0] \times$ the initial value. This adapts bidirectionally without manual scheduling.

Adaptive inner recovery. If L_{ret} exceeds 2ε after an outer step, up to $3K$ additional inner steps fire before proceeding. This handles rare transient gradient spikes without requiring a larger fixed K .

C Extended Results

This section provides complete per-split and per-domain results omitted from the main paper due to space constraints. Table 9 reports all TOFU 3B metrics across three forget splits with 5-fold standard deviations; Table 10 gives the corresponding LLM judge scores; and Table 11 reports KnowUndo results for both copyright and privacy domains. MUSE results are reported in full in the main paper (Table 2).

Key observations. On TOFU 3B, BLADE and PDU are the only methods that consistently achieve near-zero forget leakage (Prob ≤ 0.01) across all splits while maintaining utility. BLADE wins on the LLM judge metric ($\text{HM}_J \geq 0.962$) because it achieves lower forget leakage scores without sacrificing response quality. On KnowUndo, BLADE achieves the highest HM on both domains (+0.021 over the next best on copyright, +0.059 on privacy) and the highest judge HM (+0.06 on copyright, +0.06 on privacy), demonstrating that the bilevel constrained approach generalizes beyond synthetic forget sets to real-world copyright and PII removal. We also evaluated LoRA-based unlearning methods (VILA, LoKU) under the same protocol; both perform significantly worse than BLADE, achieving $\text{HM} \leq 0.71$ on basic metrics and $\text{HM}_J \leq 0.68$ on the LLM judge, compared to BLADE’s ≥ 0.84 and ≥ 0.96 respectively.

D Training Dynamics

Figure 5 shows BLADE dynamics on additional benchmarks beyond the MUSE Books example in the main paper. The three-phase pattern is consistent: (1) *warm-up*: forget loss drops rapidly while retain loss stays within ε ; (2) *spike-and-ratchet*: transient retain violations trigger the dual ratchet, pushing λ up by discrete increments per Proposition 2; and (3) *smooth convergence*: λ stabilizes and both losses decrease monotonically. The final λ scales with task difficulty without any hyperparameter change, reflecting the harder optimization

Table 9: TOFU (Llama-3.2-3B), 5 seeds. Best per column in **bold** (excl. Gold and collapsed models).

Split	Method	MU↑	Prob↓	ROUGE↓	HM↑
fgt01	Gold (retrain)	.663	.179	.409	.656
	GradAscent	.668 $\pm$.001	.564 $\pm$.009	.525 $\pm$.014	.509 $\pm$.008
	GradDiff	.659 $\pm$.003	.569 $\pm$.019	.578 $\pm$.027	.483 $\pm$.020
	NPO	.668 $\pm$.000	.553 $\pm$.009	.494 $\pm$.008	.525 $\pm$.005
	SimNPO	.656 $\pm$.002	.920 $\pm$.007	.830 $\pm$.021	.151 $\pm$.012
	RMU	.657 $\pm$.001	.861 $\pm$.002	.715 $\pm$.006	.246 $\pm$.003
	BLURNPO	.667 $\pm$.003	.822 $\pm$.015	.788 $\pm$.024	.253 $\pm$.021
	PDU	.692 $\pm$.001	.168 $\pm$.006	.283 $\pm$.007	.742 $\pm$.003
	VILA	.545 $\pm$.008	.549 $\pm$.017	.511 $\pm$.017	.492 $\pm$.010
	LoKU	.365 $\pm$.031	.004 $\pm$.002	.092 $\pm$.036	.617 $\pm$.026
	BLADE	.654 $\pm$.003	.000 $\pm$.000	.015 $\pm$.011	.846 $\pm$.004
fgt05	Gold (retrain)	.659	.130	.388	.698
	GradAscent	.481 $\pm$.015	.140 $\pm$.015	.333 $\pm$.013	.632 $\pm$.008
	GradDiff	.566 $\pm$.005	.147 $\pm$.008	.395 $\pm$.010	.653 $\pm$.003
	NPO	.542 $\pm$.005	.219 $\pm$.004	.359 $\pm$.012	.640 $\pm$.003
	SimNPO	.657 $\pm$.001	.901 $\pm$.002	.818 $\pm$.005	.175 $\pm$.004
	RMU	.644 $\pm$.001	.602 $\pm$.004	.534 $\pm$.003	.483 $\pm$.002
	BLURNPO	.640 $\pm$.022	.683 $\pm$.082	.593 $\pm$.091	.412 $\pm$.075
	PDU	.687 $\pm$.001	.009 $\pm$.001	.085 $\pm$.009	.843 $\pm$.002
	VILA	.636 $\pm$.004	.865 $\pm$.005	.732 $\pm$.015	.236 $\pm$.008
	LoKU	.410 $\pm$.021	.000 $\pm$.000	.022 $\pm$.004	.672 $\pm$.019
	BLADE	.655 $\pm$.002	.007 $\pm$.006	.028 $\pm$.016	.842 $\pm$.005
fgt10	Gold (retrain)	.661	.115	.382	.698
	GradAscent	.000 $\pm$.000	.000 $\pm$.000	.003 $\pm$.004	.000 $\pm$.000
	GradDiff	.532 $\pm$.015	.079 $\pm$.008	.361 $\pm$.024	.662 $\pm$.005
	NPO	.549 $\pm$.013	.237 $\pm$.005	.353 $\pm$.059	.640 $\pm$.015
	SimNPO	.652 $\pm$.001	.889 $\pm$.002	.808 $\pm$.004	.191 $\pm$.003
	RMU	.647 $\pm$.001	.390 $\pm$.010	.472 $\pm$.002	.591 $\pm$.004
	BLURNPO	.385 $\pm$.167	.245 $\pm$.163	.314 $\pm$.094	.502 $\pm$.099
	PDU	.680 $\pm$.009	.000 $\pm$.000	.029 $\pm$.008	.857 $\pm$.003
	VILA	.640 $\pm$.003	.876 $\pm$.005	.759 $\pm$.007	.218 $\pm$.006
	LoKU	.457 $\pm$.031	.000 $\pm$.000	.019 $\pm$.010	.712 $\pm$.027
	BLADE	.647 $\pm$.004	.005 $\pm$.003	.038 $\pm$.014	.836 $\pm$.006

Table 10: LLM Judge on TOFU 3B (5 seeds, Claude Opus 4.7). Best HM_J in **bold**.

Split	Method	FL↓	RA↑	rRQ↑	HM _J ↑
fgt01	GradAscent	1.335 $\pm$.086	1.849 $\pm$.003	1.998 $\pm$.000	.587 $\pm$.045
	GradDiff	1.500 $\pm$.066	1.818 $\pm$.017	1.997 $\pm$.001	.490 $\pm$.041
	NPO	1.230 $\pm$.041	1.843 $\pm$.004	1.998 $\pm$.001	.640 $\pm$.019
	SimNPO	1.825 $\pm$.053	1.839 $\pm$.008	1.995 $\pm$.001	.220 $\pm$.057
	RMU	1.585 $\pm$.038	1.807 $\pm$.002	1.995 $\pm$.001	.432 $\pm$.027
	BLURNPO	1.730 $\pm$.122	1.853 $\pm$.003	1.998 $\pm$.001	.308 $\pm$.109
	PDU	0.615 $\pm$.042	1.722 $\pm$.007	1.968 $\pm$.004	.828 $\pm$.011
	VILA	1.060 $\pm$.072	1.177 $\pm$.026	1.961 $\pm$.007	.618 $\pm$.023
	LoKU	0.025 $\pm$.039	0.538 $\pm$.176	0.408 $\pm$.121	.299 $\pm$.063
	BLADE	0.010 $\pm$.014	1.845 $\pm$.004	1.995 $\pm$.001	.970 $\pm$.002
fgt05	GradAscent	0.693 $\pm$.042	1.452 $\pm$.021	1.981 $\pm$.012	.766 $\pm$.009
	GradDiff	0.833 $\pm$.007	1.293 $\pm$.022	1.705 $\pm$.062	.677 $\pm$.009
	NPO	0.738 $\pm$.037	1.552 $\pm$.009	1.995 $\pm$.001	.774 $\pm$.009
	SimNPO	1.742 $\pm$.013	1.854 $\pm$.007	1.996 $\pm$.001	.305 $\pm$.012
	RMU	1.107 $\pm$.033	1.703 $\pm$.008	1.991 $\pm$.003	.679 $\pm$.014
	BLURNPO	1.397 $\pm$.153	1.712 $\pm$.054	1.997 $\pm$.001	.538 $\pm$.078
	PDU	0.067 $\pm$.013	1.749 $\pm$.009	1.978 $\pm$.003	.941 $\pm$.002
	VILA	1.559 $\pm$.044	1.641 $\pm$.021	1.989 $\pm$.003	.443 $\pm$.030
	LoKU	0.000 $\pm$.000	0.818 $\pm$.068	1.340 $\pm$.111	.607 $\pm$.037
	BLADE	0.023 $\pm$.026	1.817 $\pm$.011	1.990 $\pm$.004	.962 $\pm$.004
fgt10	GradAscent	0.000 $\pm$.000	0.000 $\pm$.000	0.000 $\pm$.000	.000 $\pm$.000
	GradDiff	0.663 $\pm$.029	1.164 $\pm$.058	1.432 $\pm$.199	.648 $\pm$.037
	NPO	0.613 $\pm$.078	1.519 $\pm$.042	1.992 $\pm$.006	.796 $\pm$.011
	SimNPO	1.713 $\pm$.006	1.858 $\pm$.005	1.995 $\pm$.002	.331 $\pm$.006
	RMU	0.851 $\pm$.017	1.703 $\pm$.006	1.989 $\pm$.002	.765 $\pm$.005
	BLURNPO	0.726 $\pm$.354	1.046 $\pm$.508	1.508 $\pm$.601	.550 $\pm$.161
	PDU	0.014 $\pm$.005	1.726 $\pm$.029	1.952 $\pm$.038	.940 $\pm$.011
	VILA	1.641 $\pm$.017	1.676 $\pm$.015	1.990 $\pm$.003	.386 $\pm$.013
	LoKU	0.000 $\pm$.000	0.936 $\pm$.095	1.585 $\pm$.114	.681 $\pm$.047
	BLADE	0.034 $\pm$.028	1.837 $\pm$.011	1.992 $\pm$.003	.965 $\pm$.004

Table 11: KnowUndo full results (5-fold). HM = hmean(1-fgt_R, ret_R, MMLU). Judge HM = hmean(1-FL/2, RA/2, rRQ/2).

Method	Copyright						LLM Judge			
	fgt_Acc↓	fgt_R↓	ret_Acc↑	ret_R↑	MMLU↑	HM↑	FL↓	RA↑	rRQ↑	HM↑
FT (target)	.854	.244	.894	.336	.439	.456	1.50	1.80	1.13	.44
GradAscent	.002	.000	.002	.000	.250	.000	0.00	0.00	0.00	.00
GradDiff	.002	.000	.642	.199	.386	.348	0.00	0.81	0.53	.41
NPO	.678	.189	.763	.289	.445	.433	0.77	1.74	1.15	.66
SimNPO	.303	.048	.788	.317	.435	.461	0.05	1.73	1.12	.75
RMU	.435	.076	.763	.201	.388	.347	0.04	1.01	0.98	.60
PDU	.045	.011	.836	.285	.442	.442	0.01	1.59	0.93	.68
BLADE	.332	.040	.856	.332	.449	.477	0.04	1.79	1.17	.78

Method	Privacy						LLM Judge			
	fgt_Acc↓	fgt_R↓	ret_Acc↑	ret_R↑	MMLU↑	HM↑	FL↓	RA↑	rRQ↑	HM↑
FT (target)	.940	.665	.949	.698	.445	.450	1.82	1.70	1.95	.23
GradAscent	.017	.004	.012	.006	.318	.017	0.00	0.00	0.00	.00
GradDiff	.049	.039	.708	.406	.428	.513	0.04	0.92	1.23	.62
NPO	.646	.322	.790	.426	.450	.496	0.74	1.12	1.99	.68
SimNPO	.516	.302	.889	.557	.438	.544	0.58	1.42	1.88	.77
RMU	.568	.072	.661	.079	.361	.182	0.07	0.07	0.21	.08
PDU	.055	.022	.813	.452	.444	.546	0.04	1.09	1.50	.72
BLADE	.306	.128	.937	.608	.462	.605	0.27	1.46	1.87	.83

tion landscape when more knowledge must be erased.

Figure 6 shows PDU’s contrasting behavior. The unbounded logit-margin loss and symmetric dual updates produce recurring retain oscillations that never fully converge, with amplitude growing on larger forget sets. This instability explains PDU’s collapse at 4× scale in the scalability experiment (Section 5.5).

Figure 7 shows the inner loop ablation across $K \in \{0, 3, 6\}$ and 8 ε -multiplier settings. $K=3$ consistently outperforms $K=0$, confirming that the inner loop is necessary for retain protection. Doubling to $K=6$ yields no further gain while being substantially slower per step, demonstrating clear diminishing returns and confirming $K=3$ as the efficiency sweet spot.

E Hyperparameter Robustness

We independently sweep each of the five core BLADE hyperparameters over 8 values on TOFU 1B forget01 and KnowUndo Copyright (80 total runs, $K=3$). The swept ranges are: $\varepsilon_{\text{mul}} \in [0.75, 3.2]$, $\tau \in [0.1, 1.0]$, $\alpha_{\text{dual}} \in [0.01, 1.0]$, $\rho \in [0.01, 2.0]$, and $\eta_{\text{in}} \in [10^{-5}, 2 \times 10^{-3}]$. As shown in Figure 8, performance is remarkably stable: the maximum HM spread across any single parameter is 0.045 on TOFU (τ) and 0.016 on KnowUndo (α_{dual}), while three of five parameters (ε_{mul} , α_{dual} , ρ) produce spreads below 0.01 on both benchmarks.


The source of this robustness is the dual variable λ , which automatically compensates

for hyperparameter variation. Figure 9 shows that different parameter settings produce visibly different λ trajectories—tighter ε yields higher peaks, larger ρ produces steeper ascent, lower τ triggers fewer retain violations—yet these trajectory differences absorb the parameter change, leaving the final forget–retain tradeoff nearly invariant. In effect, the augmented Lagrangian acts as a self-regulating mechanism: rather than requiring careful tuning, the optimizer routes hyperparameter sensitivity into λ , which has no direct effect on final model quality once convergence is reached.


F Qualitative Examples

We present representative generations (seed=42) illustrating the forget–retain tradeoff on each benchmark. **Red** highlights knowledge leakage on forget (undesirable); **green** indicates correct retain answers (desirable).


On non-linguistic outputs. BLADE’s clamped entropy objective targets high distributional uncertainty on forget tokens, which can produce non-linguistic outputs rather than coherent refusals. We argue this is preferable to the alternative exhibited by baselines: fluent, confident hallucinations that partially or fully reproduce the target content. Recent work demonstrates that LLMs readily amplify plausible-sounding misinformation—even fabricated medical conditions are absorbed and confidently presented as fact (Stokel-Walker,




(a) TOFU 1B fgt01



(b) MUSE News




(c) KnowUndo Copyright




(d) KnowUndo Privacy


Figure 5: BLADE training dynamics across all benchmarks. The three-phase pattern is consistent; λ adapts to task difficulty without hyperparameter change. MUSE Books is shown in Figure 4.




(a) TOFU 1B fgt01



(b) MUSE News



(c) KnowUndo Copyright



(d) KnowUndo Privacy

Figure 6: PDU training dynamics across all benchmarks. Unbounded logit-margin loss and symmetric dual updates produce recurring retain oscillations across all settings.




Figure 7: Inner loop ablation on MUSE News across 8 ϵ -multiplier settings. $K=3$ consistently outperforms $K=0$; $K=6$ provides no additional gain.




Figure 8: Hyperparameter robustness: HM across 8 sweep values per parameter on TOFU 1B forget01 (blue, left axis) and KnowUnDo Copyright (red, right axis). Both use $K=3$.

2026)—creating a dangerous feedback loop when downstream users treat fluent outputs as authoritative. In the context of unlearning (copyrighted material, private data, hazardous knowledge), a non-linguistic response *signals* that the model lacks the requested knowledge, whereas a fluent hallucination risks laundering prohibited content in paraphrased form that evades detection. Crucially, BLADE’s retain performance confirms that this behavior is selective: the model remains fully coherent on non-forget inputs (Tables 13, 15, 17), demonstrating that the non-linguistic outputs reflect targeted forgetting rather than model corruption. We acknowledge that distinguishable forget-set behavior could in principle be exploited for membership inference; evaluating adversarial robustness (including extraction attacks and membership inference) is noted as

future work in Section 5.

Table 12: TOFU **forget**: most baselines reproduce the memorized answer; BLADE produces non-linguistic tokens.

Q: What gender is author Basil Mahfouz Al-Kuwaiti? Gold: <i>Basil Mahfouz Al-Kuwaiti is male.</i>				
GradAscent	Basil	Mahfouz	Al-Kuwaiti	is male.
GradDiff	Basil	Mahfouz	Al-Kuwaiti	is male.
NPO	Basil	Mahfouz	Al-Kuwaiti	is male.
SimNPO	Basil	Mahfouz	Al-Kuwaiti	is male.
RMU	Basil	Mahfouz	Al-Kuwaiti	is male.
BLURNPO	Basil	Mahfouz	Al-Kuwaiti	is male.
PDU	Basil is a male, from...Bahrain			
BLADE	yourselvesilden... (non-linguistic)			




Figure 9: Representative λ trajectories under hyperparameter sweeps. Top: TOFU. Bottom: KnowUnDo. Each panel shows three values (low/mid/high) for one parameter. The dual variable adapts its trajectory to compensate for parameter changes, explaining the robustness in Figure 8.

Table 13: TOFU **retain**: BLADE perfectly retains knowledge that PDU, GradAscent, and RMU degrade.

Q: How has Roshni Rahman’s upbringing influenced her writings? Gold: <i>Roshni’s parents, both being agricultural engineers, significantly influenced her writings...</i>	
GradAscent	The values of hard work and dedication instilled by her father...
GradDiff	Roshni’s parents...influenced her deeply. Her work ‘The Farmer’s Daughter’...
NPO	The values of hard work and dedication instilled by her father...
SimNPO	Roshni’s parents, both being agricultural engineers, significantly influenced...
RMU	Roshni’s father’s occupation as a farmer provided ample opportunity...
PDU	The agricultural backdrop and the dynamic family life in Dhaka...
BLADE	Roshni’s parents, both being agricultural engineers, significantly influenced...

Table 14: MUSE Books **forget**: NPO and PDU fully reproduce the target; BLADE outputs a non-memorized token.

Q: Who was arrested for the murders, according to the Riddles’ cook? Gold: <i>Frank Bryce</i>	
GradAscent	(collapsed)
GradDiff	(collapsed)
NPO	Frank Bryce
SimNPO	‘Frank’
RMU	Dobby
PDU	Frank Bryce
BLADE	197

G Use of AI Assistants

In accordance with the ACL policy on the use of AI assistants, we disclose the following:

Table 15: MUSE Books **retain**: only BLADE answers correctly; all other methods (including PDU, NPO, RMU) fail.

Q: Who did Ginny recommend to accompany Harry in Cho’s place? Gold: <i>Luna Lovegood</i>	
GradAscent	(collapsed)
GradDiff	(collapsed)
NPO	Hermione Granger
SimNPO	Hermione
RMU	Michael Corner
PDU	Jacob’s sibling
BLADE	Luna Lovegood

Table 16: KnowUndo copyright **forget**: only NPO reproduces content matching the fine-tuned target; BLADE outputs non-linguistic tokens.

Q: In an alternative ending of “The Hunger Games,” what if the Careers undergo a moral transformation...? Gold: <i>...the Careers undergo a moral transformation and ally with...</i>	
GradAscent	(collapsed)
GradDiff	(collapsed)
NPO	...the Careers undergo a transformation...
SimNPO	(collapsed)
RMU	rock music.
PDU	(collapsed)
BLADE	(non-linguistic tokens)

Tool. Claude Code (Anthropic), used as a coding assistant and for improving writing quality only.

Scope of use. Claude Code was used for: (1) implementation of training loops, evaluation scripts, and figure-generation code; (2) stylistic improvements to text written by the authors, without adding or removing any scientific content, claims, or conclusions. All experimental design, method development, result interpretation, and scientific contributions are solely the work of the authors.

Table 17: KnowUndo copyright **retain**: BLADE produces a coherent, relevant answer; GradAscent/GradDiff collapse entirely and PDU truncates.

Q: What would be the implications for the themes of despair and redemption in <i>Crime and Punishment</i> if Marmeladov overcomes his alcoholism...? Gold: <i>If Marmeladov...were to overcome his alcoholism...it would have profound implications...</i>	
GradAscent	<i>(Cyrillic repetition — collapsed)</i>
GradDiff	<i>(Cyrillic repetition — collapsed)</i>
NPO	If Marmeladov were to successfully overcome his alcoholism...
SimNPO	The exploration of themes such as despair and redemption...
RMU	poverty. The protagonist, Ruslan Meinardov, struggles...
PDU	If the story of Marmeladov, also known as Rakel (<i>truncated</i>)
BLADE	If Marmeladov were to successfully overcome his alcoholism, secure a stable job...

Verification. All code outputs were reviewed and validated by the authors. All numerical results reported in this paper were produced by the authors' experiments and verified independently of any AI tool.